



**CYBER
CHALLENGE**
CyberChallenge.it



SPONSOR PLATINUM

accenture security

aizoon AUSTRALIA
EUROPE USA
TECHNOLOGY CONSULTING

B5

EY Building a better
working world



expravia | **ITALTEL**

IBM

KPMG

LEONARDO

NTT DATA
Trusted Global Innovator

NUMERA
SISTEMI E INFORMATICA S.p.A.

Telsy

SPONSOR GOLD

bip.

CISCO

**MONTE
DEI PASCHI
DI SIENA**
BANCA DAL 1472

negg®

NOVANEXT
connecting the future

pwc

SPONSOR SILVER

**DGi
ONE**
the leading
digital company

**ICT
CYBER
CONSULTING**

Software Security 2

Memory Management and Allocation

2

Michele LORETI

Univ. di Camerino

michele.loreti@unicam.it



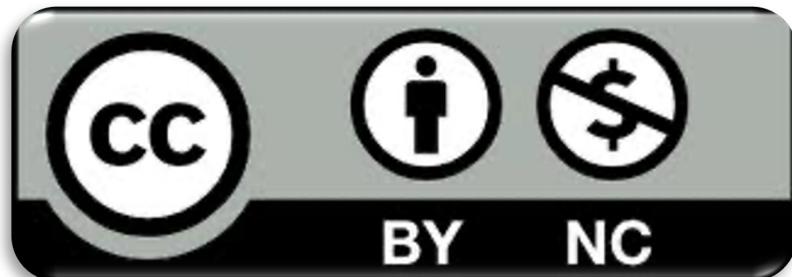
<https://cybersecnatlab.it>

License & Disclaimer

3

License Information

This presentation is licensed under the Creative Commons BY-NC License



To view a copy of the license, visit:

<http://creativecommons.org/licenses/by-nc/3.0/legalcode>

Disclaimer

- We disclaim any warranties or representations as to the accuracy or completeness of this material.
- Materials are provided “as is” without warranty of any kind, either express or implied, including without limitation, warranties of merchantability, fitness for a particular purpose, and non-infringement.
- Under no circumstances shall we be liable for any loss, damage, liability or expense incurred or suffered which is claimed to have resulted from use of this material.

Outline

4

- Memory Management
- Memory Layout
- Stack
- Heap
- Bad stories

Memory management

5

- Many of the modern programming languages allow programmers to use **data types** without having to know **how** they are represented.
- Similarly, programmers often ignore **where** data is stored (**allocated**)...
 - compiler makes decisions, however at runtime allocation is under the control of Operating System and CPU.

Memory management

6

- In some programming languages, like C, memory management can be controlled by programmers:
 - memory can be **dynamically** allocated and deallocated;
 - memory address of variables can be obtained (pointers).
- If x is a variable, $\&x$ denotes **the pointer to x** , i.e., the memory address where x is stored.

Memory allocation...

7

Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

We can assume that:

- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes.

Memory allocation...

8

Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l; } } Variable declarations

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l); }
```

We can assume that:

- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes.

Memory allocation...

9

Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;      Variable addresses

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

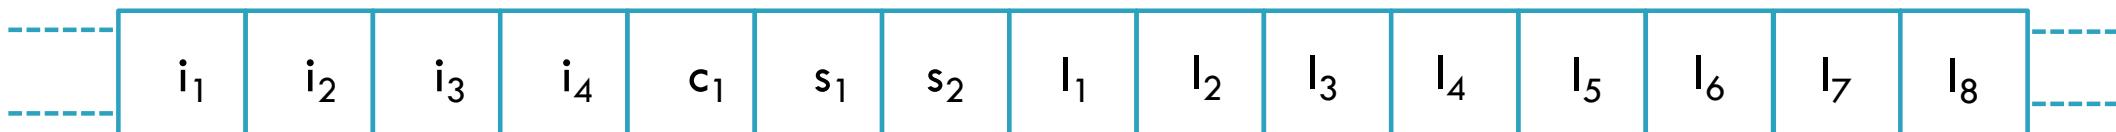
We can assume that:

- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes.

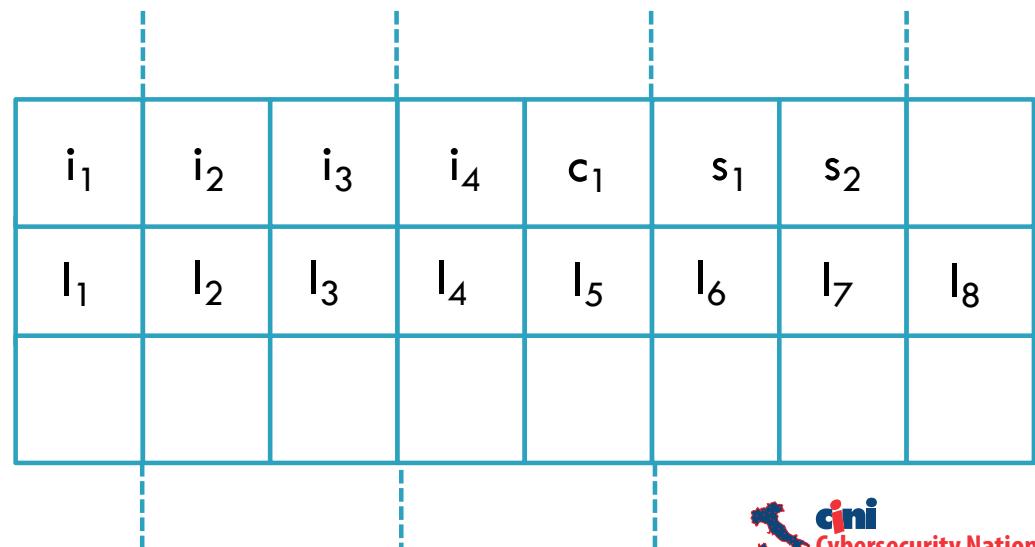
Memory allocation...

10

Memory is usually represented as a sequence of bytes:



In a n·8 architecture, bytes are arranged in groups of n:



Memory allocation...

11

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

We can run this simple program to observe how data are allocated in memory.

Memory allocation...

12

Different allocation policies can be used by *gcc*:

```
[CC> gcc -o alignment alignment.c
[CC> ./alignment
i is allocated at 0x7ffee117e7cc
c is allocated at 0x7ffee117e7cb
s is allocated at 0x7ffee117e7c8
l is allocated at 0x7ffee117e7c0
[CC>
```

```
[CC> gcc -O2 -o alignment alignment.c
[CC> ./alignment
i is allocated at 0x7ffee4dbb7c8
c is allocated at 0x7ffee4dbb7cf
s is allocated at 0x7ffee4dbb7cc
l is allocated at 0x7ffee4dbb7c0
[CC>
```

Data alignment

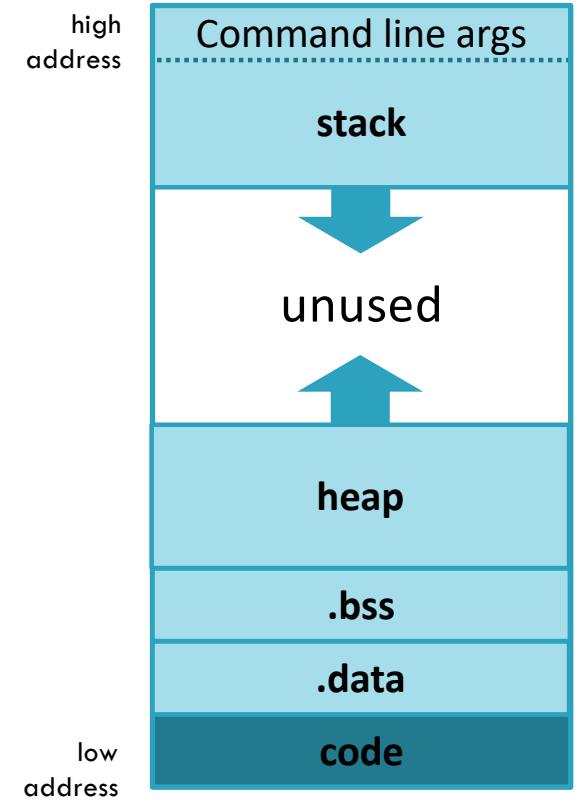
13

- Compilers may introduce **padding** or change the order of data in memory
- There are trade-offs between speed and memory usage.
- C compilers provide many optional optimization.

Memory segments

14

- Memory is allocated for each **process** (a running program) to store **data** and **code**.
- This allocated memory consists of different **segments**:
 - **stack**: for local variables
 - **heap**: for dynamic memory
 - **data segment**:
 - *global uninitialized variables (.bss)*
 - *global initialized variables (.data)*
 - **code segment**



The stack

15

- The stack consists of a sequence of stack frames (or activation records), each for each function call:
 - allocated on *call*;
 - de-allocated on *return*.

```
int main(int argc, char **argv) {  
    int f = fib( n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib( n: n-1);  
        f2 = fib( n: n-2);  
        return f1+f2;  
    }  
}
```

stack frame
for main()

stack frame
for fib()

stack frame
for fib()

Unused memory

The stack

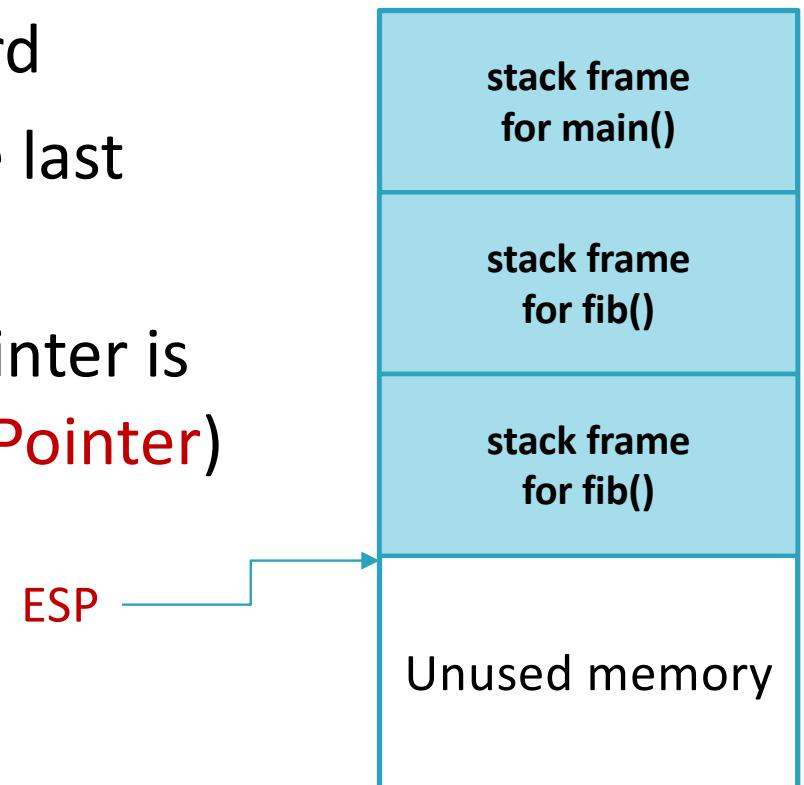
16

- The precise structure and organization of the stack depends on system architecture, operating system, and compilers we are using.
- For the sake of simplicity, in this lecture we will focus on x86 architectures (32 and 64 bits) and on *gcc* compiler on *Linux*.
- More detailed information are available at the following links:
 - <https://docs.microsoft.com/en-us/cpp/cpp/argument-passing-and-naming-conventions>
 - <http://refspecs.linuxbase.org/>

The stack

17

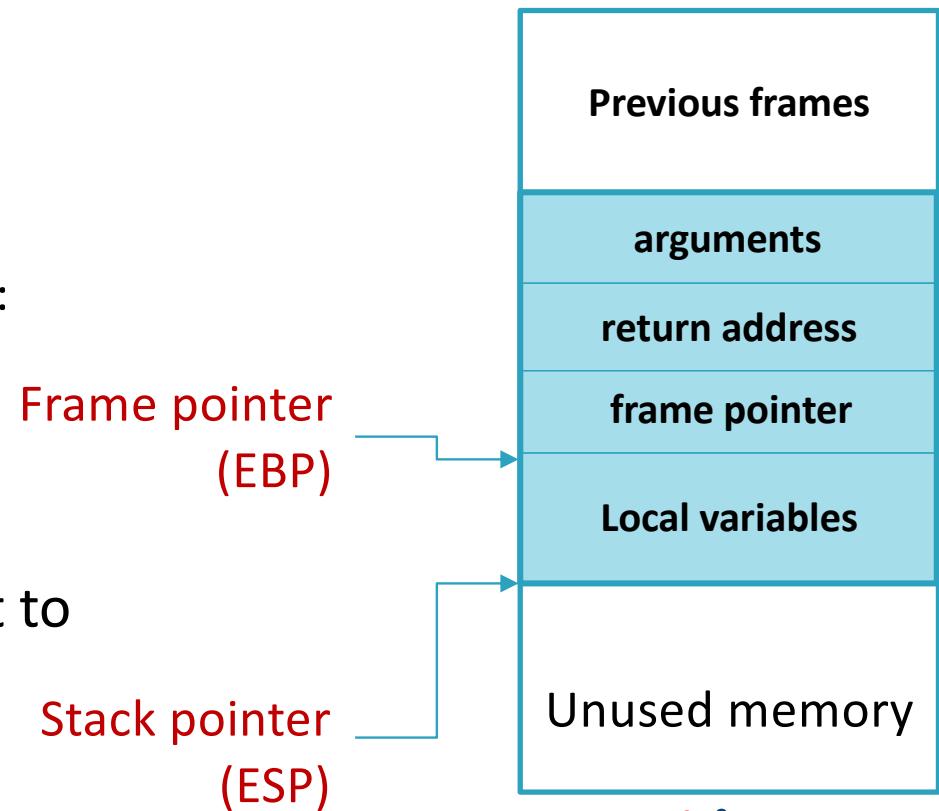
- Typically the stack grows downward
- The **stack pointer (SP)** refers to the last element on the stack
- On x86 architectures, the stack pointer is stored in the **ESP (Extended Stack Pointer)** register.



Stack frame (for x86)

18

- In x86 architecture, each stack frame contains:
 - Function arguments;
 - Local variables;
 - Copies of registries that must be restored:
 - return address
 - previous frame pointer
- Frame pointer, named Extended Base Pointer (EBP), provides a starting point to local variables.



Stack frame: example

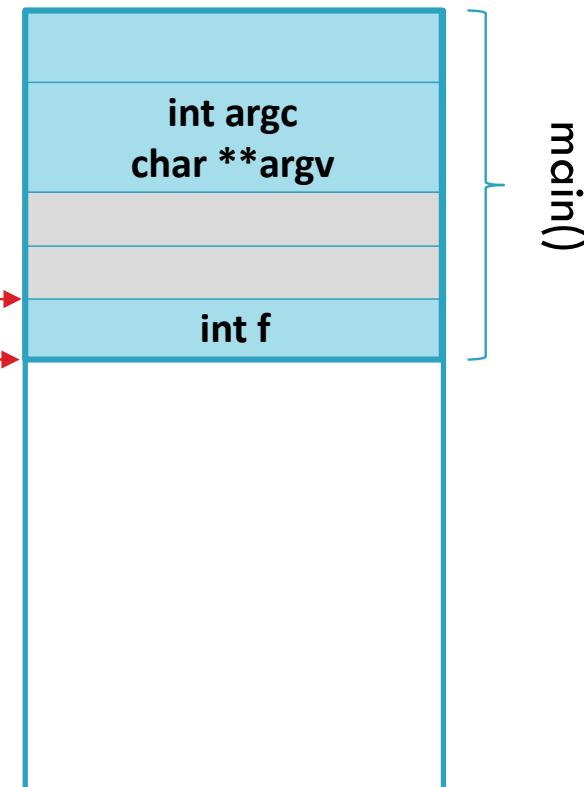
19

```
int main(int argc, char **argv) {
    int f = fib( n: 10); ←
    printf("FIB(10)=%d\n", f);
}

int fib(int n) {
    int f1;
    int f2;
    if (n<=2) {
        return 1;
    } else {
        f1 = fib( n: n-1);
        f2 = fib( n: n-2);
        return f1+f2;
    }
}
```

Frame pointer
Stack pointer

Function fib is
invoked with
parameter 10.



Stack frame: example

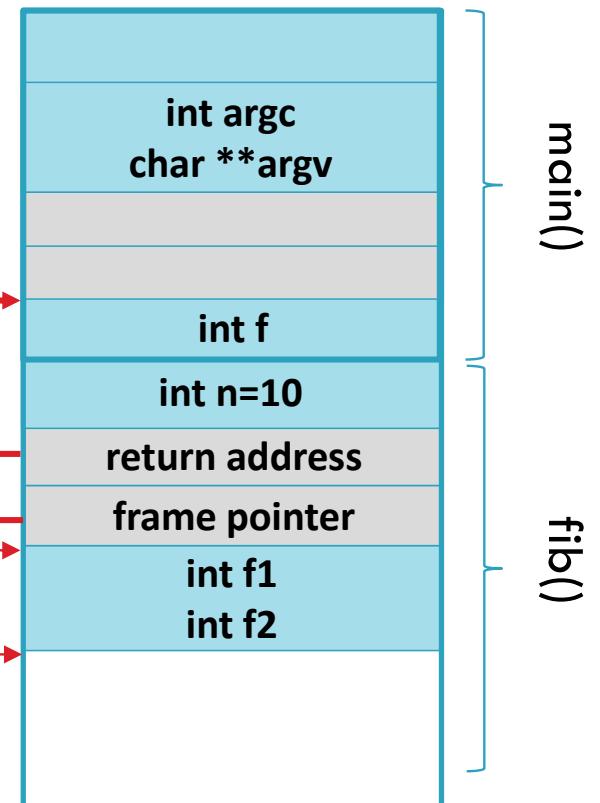
20

```
int main(int argc, char **argv) {  
    int f = fib( n: 10);  
    printf("FIB(10)=%d\n",f);  
}  
  
int fib(int n) {  
    int f1; ←  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib( n: n-1);  
        f2 = fib( n: n-2);  
        return f1+f2;  
    }  
}
```

Stack frame is allocated and pointers updated.

Frame pointer

Stack pointer



Stack frame: example

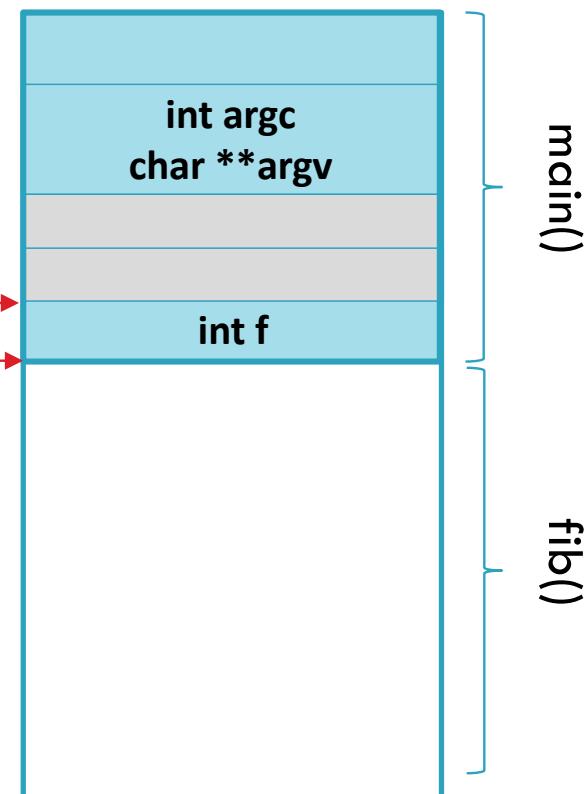
21

```
int main(int argc, char **argv) {
    int f = fib( n: 10); ←
    printf("FIB(10)=%d\n",f);
}

int fib(int n) {
    int f1;
    int f2;
    if (n<=2) {
        return 1;
    } else {
        f1 = fib( n: n-1);
        f2 = fib( n: n-2);
        return f1+f2;
    }
}
```

Frame pointer
Stack pointer

When a function returns, pointers are updated. Function result (if any) is copied in a register.



Stack frame (for x86-64)

22

- In the *64-bit* version of *x86*:
 - Registers are extended to 64 bits;
 - A new set of 8 registers are added.
- The arguments are not all in the stack:
 - The first 6 are passed via registers;
 - The remaining are placed in the stack (like for *x86*).
- Pointers *EBP* and *ESP* are named *RBP* and *RSP*.
- A *red zone* of 128 bytes is placed in the stack just under RSP
 - This can be used to store extra local variables and is not modified by interrupt/exception/signal handlers.

C declaration: cdecl

23

- The **cdecl** (C declaration) is a **calling convention** used in many compilers for *x86*
 - A calling convention describes how functions receive their parameters from caller and how they return their results.
- In cdecl arguments are passed on the stack, while values are returned via registers.
- Function arguments are pushed on the stack in the *right-to-left order*.
- The caller *cleans* the stack after the function call returns.

Security issues

24

- In the stack, variables do not have any **default initialization**:
 - uninitialized local variables can be used to read values from previous function calls;
- The space assigned to the stack is finite and a function call may fail because there is no more memory:
 - We must guarantee that this does not happen or handle it in a safe way.
- Program and control data are mixed in the stack:
 - By accessing over buffers limit may corrupt the return address!

The heap

25

- Memory allocation and de-allocation in the stack is very fast
 - However, this memory cannot be used after a function returns
- The heap is used to store dynamically allocated data that outlive function calls:
 - This area is under programmer's responsibility.

Memory management functions

26

- Basic C functions for memory management are:
 - *malloc(int)*, given an integer n allocates an area of n (continuous) bytes and returns a **pointer** to that area;
 - *free(void*)*, deallocates the memory associated with a pointer.

Stack and Heap pointers

27

- We have always to preserve reference to allocated areas in the heap
 - Otherwise we cannot use them!
- These references (pointers) can be stored in the stack or in the heap.
- We could store pointers to the stack in the heap
 - This can be dangerous! Referenced memory could be released!

Memory security problems

28

- In C memory safety is not guaranteed and malicious or insecure code can access data anywhere on stack and heap:
 - via pointer arithmetic;
 - by overrunning array bounds.
- Security problems with memory are due to:
 - Running out of memory;
 - Lack of initialization of memory;
 - Bugs in program code.

List of common errors (1/3)

29

- Dereferencing an uninitialized or freed address:

```
int *p;  
int z = *p;
```

```
int *p = malloc(sizeof(int));  
...  
free(p)  
...  
int z = *p;
```

In both the cases we are going to use an invalid memory location. The result is unpredictable.

List of common errors (2/3)

30

- Freeing an already freed pointer or a pointer that was not dynamically allocated :

```
int *p = malloc(sizeof(int));
...
free(p)
...
free(p)
```

```
int z = 10;
int *p = &z;
...
free(p)
```

In both the cases the result is unpredictable. It may cause abnormal termination, or memory corruption.

List of common errors (3/3)

31

- Attempt to access to memory beyond the allocated buffer:

```
int *p1 = malloc(100*sizeof(int));
int *p2 = malloc(200*sizeof(int));

p2=p1;

p2[150]=10;
```

We have an overflow: a memory area outside the buffer is changed!

Good practices

32

- Every *malloc* should have an associated *free*.
- Pointers should be initialized when defined (either to a valid address or to *NULL*).
- Pointers should be assigned to *NULL* after being freed.
- Test if a pointer is *NULL* before using it (**defensive programming**).

Bad stories...

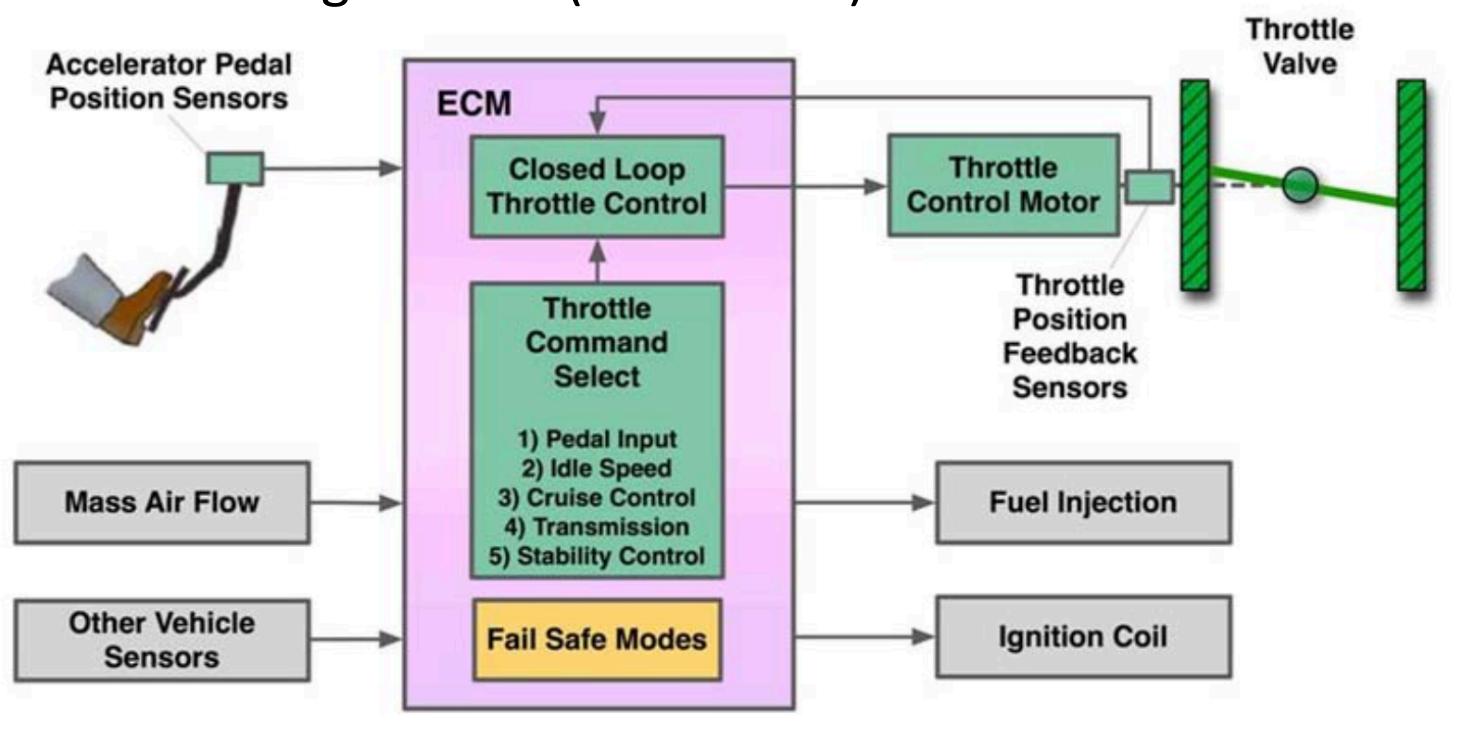
33

- Bugs in memory allocation can be the source of serious software vulnerabilities.
- Two well known examples are:
 - Toyota unintended acceleration (UA) cases;
 - Heartbleed bug in OpenSSL.

Unintentionally Acceleration

34

- NASA team investigates UA (2010-2011)



Unintentionally Acceleration

35

- The bug originated from a **stack overflow** that may corrupt critical variables of the Operating System.

Heartbleed bug

36

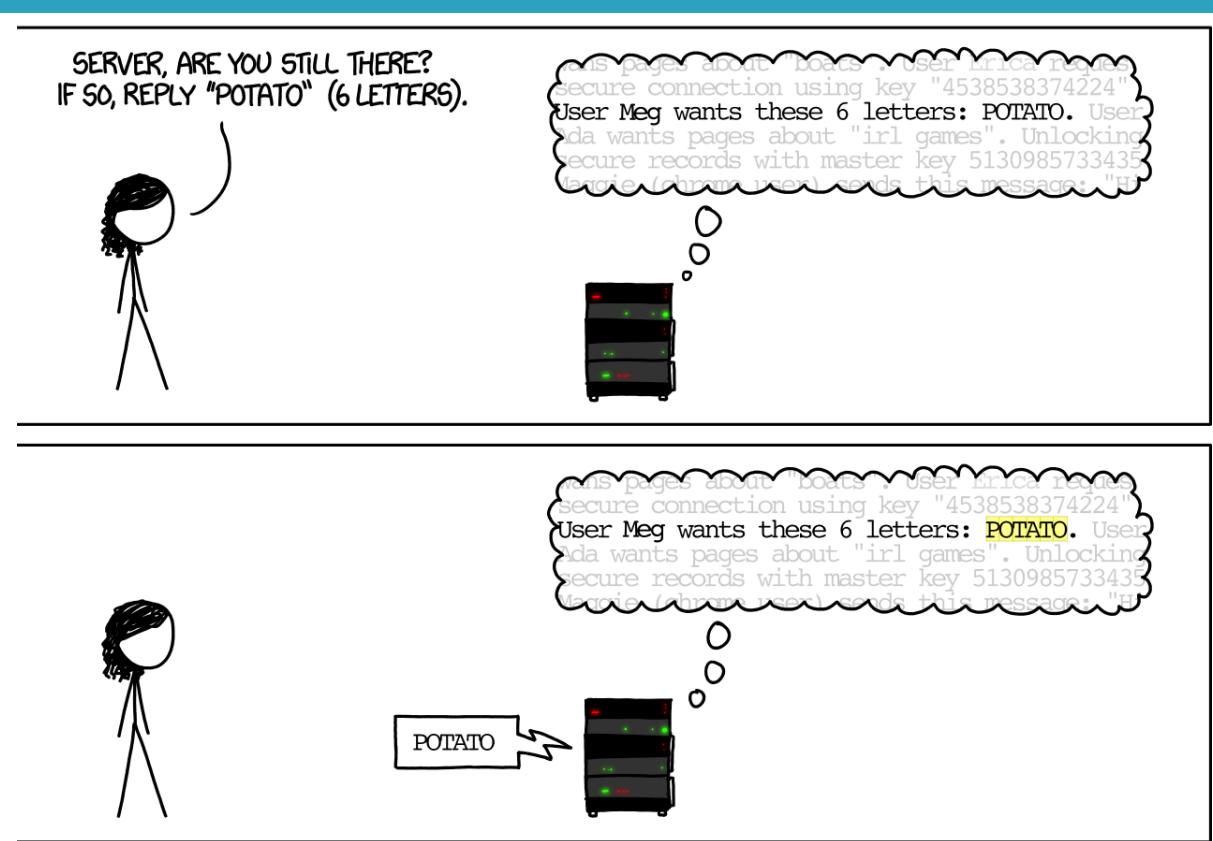
- The heartbeat functionality in SSL can be used to check if a connection is still alive.
- Malformed input can be provided to OpenSSL to let it print a large part of the stack (possibly containing private keys).
- The same problem effects some TLS implementations (see [CVE-2016-9244](#)).

Heartbleed bug

37

Xkcd provides a nice explanation how the bug works:

<https://xkcd.com/1354/>

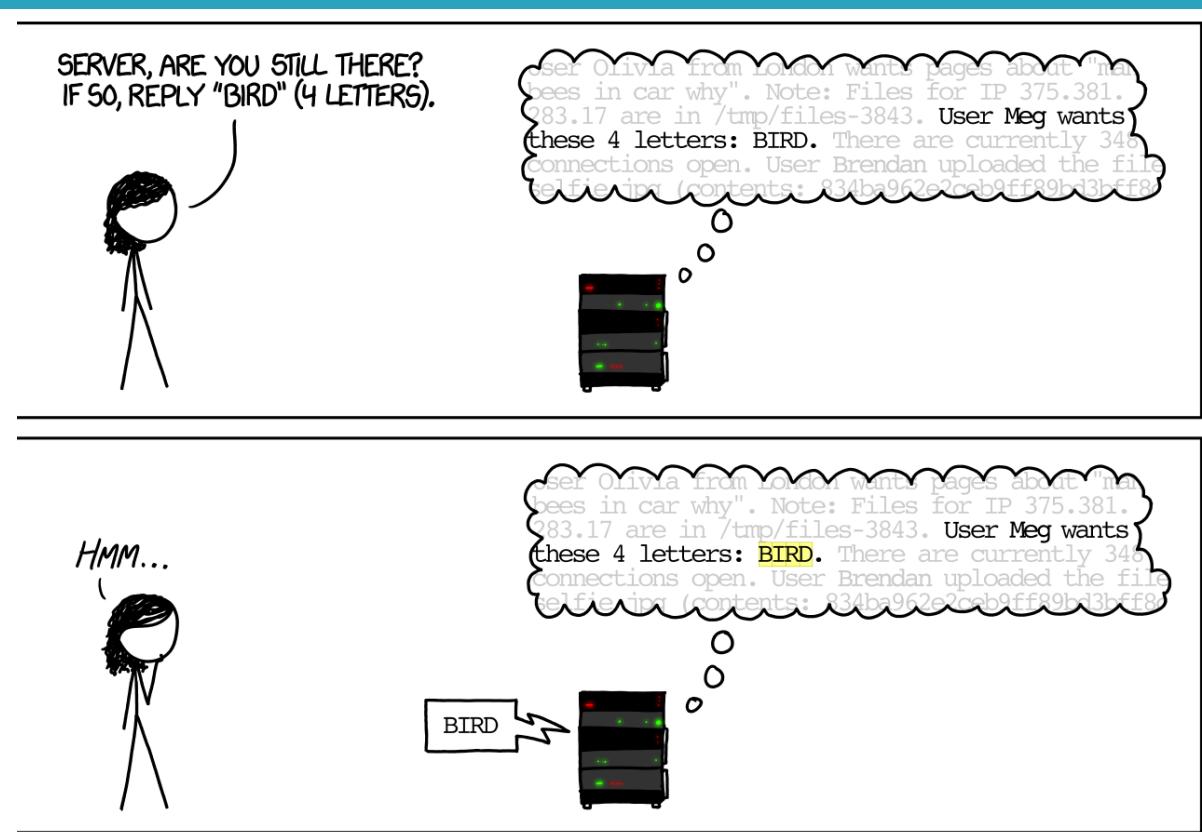


Heartbleed bug

38

Xkcd provides a nice explanation how the bug works:

<https://xkcd.com/1354/>

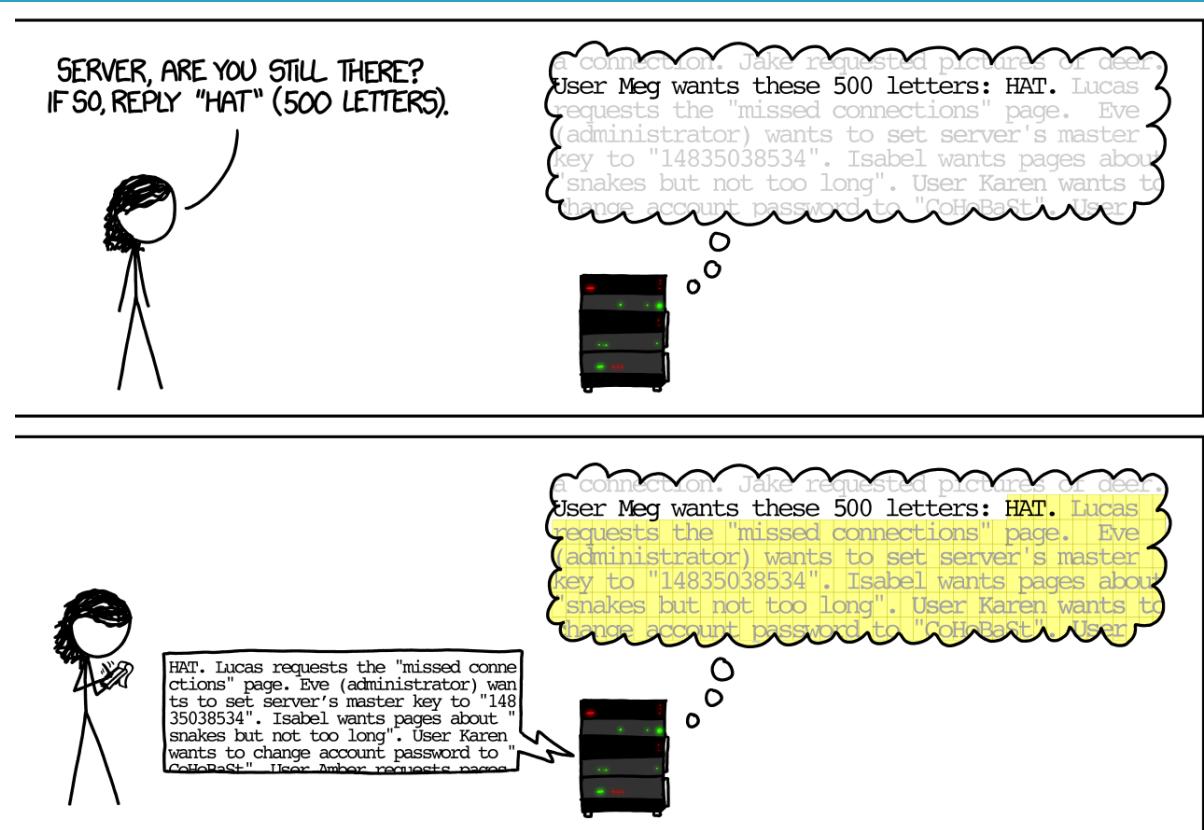


Heartbleed bug

39

Xkcd provides a nice explanation how the bug works:

<https://xkcd.com/1354/>



Software Security 2

Memory Management and Allocation

40

Michele LORETI

Univ. di Camerino

michele.loreti@unicam.it



<https://cybersecnatlab.it>



**CYBER
CHALLENGE**
CyberChallenge.it



SPONSOR PLATINUM

accenture security

aizoon AUSTRALIA
EUROPE USA
TECHNOLOGY CONSULTING

B5

EY Building a better
working world

eni

expravia | **ITALTEL**

IBM

KPMG

LEONARDO

NTT DATA
Trusted Global Innovator

NUMERA
SISTEMI E INFORMATICA S.p.A.

Telsy

SPONSOR GOLD

bip.

CISCO

**MONTE
DEI PASCHI
DI SIENA**
BANCA DAL 1472

negg®

NOVANEXT
connecting the future

pwc

SPONSOR SILVER

**DGi
ONE**
the leading
digital company

**ICT
CYBER
CONSULTING**